

1.1

steps	States (s)	Action (a)	Next State (s')	Reward (R)	Total
1	0	2	2	0.0	0.0
2	2	3	3	3.0	3.0
3	3	2	2	0.0	3.0
4	2	3	4	3.0	6.0
5	3	2	2	0.0	6.0

Highest rewarding transition belongs to the step where from state 2 action 3 was taken and next state was reached. After starting from state 0, 1 step makes it to state

2. Then we can perform state 2 to state 3 transition twice while coming back 1 time and spending 4 steps of a single episode. Last transition is from state 3 to state 2 again because rest of the other options incur negative rewards.

1.3.b

I believe the optimal achievable reward was achieved. The scores of “DejaVu Sans” are included along side score.png

```
findfont: score(<Font 'DejaVu Sans' (DejaVuSans-Bold.ttf) normal normal 700 normal>) = 0.33499999999999996
```

:

:

:

```
findfont: Matching sans\serif:style=normal:variant=normal:weight=normal:stretch=normal:size=10.0 to  
DejaVu Sans ('C:\\Users\\Nayeem\\anaconda3\\envs\\cs234-torch\\lib\\site-  
packages\\matplotlib\\mpl-data\\fonts\\ttf\\DejaVuSans.ttf') with score of 0.050000.
```




Fig1: the “scores.png” of q3_linear_torch.py

1.3.d.




Fig2: the “scores.png” of q4_nature_torch.py




Fig3: the “scores.png” of q4_nature_torch.py after config of q4_nature.py was modified. I changed the nsteps_train from 1,000 to 10,000. Time took approximately 6x of (nsteps_train=1,000)’runtime. ~5/6 minutes.


In linear approximation, it reached the optimal reward at 2nd epoch whereas in non-linear approximation it took around 5th epoch to reach optimal score of 6.2 (approximate). In both cases, the runtime was around same which is around 1 minute and both closed near the optimal reward.

When I changed the nsteps_train from 1k to 10k, it took approximately 6 times more time to train.

It takes a little bit more time for non-linear network to reach optimal reward than linear network. In this case, linear network performed slightly better than non-linear network.

1.4.

At first glance, the agent's performance steadily decreases for almost half of the training cycle and then it starts to rapidly increase for the rest of the training cycle. It decreased from -20.93 to -21.7 approximately and then rose up to -19.8 approximately. Since there wasn't any plateau which was observed, it is safe to predict that with a large number of training steps further improvements can be made.



```

findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/STIXTwoExtra.ttf", name="STIXTwoExtra", style="normal", variant="normal", weight="normal", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/cmyu101.ttf", name="cmyu101", style="normal", variant="normal", weight="normal", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/cw111.ttf", name="cw111", style="normal", variant="normal", weight="normal", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/DeJaVuSerif.ttf", name="DeJaVuSerif", style="normal", variant="normal", weight="normal", stretch="normal", size="scalable")) = 10.335
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/cw111.ttf", name="cw111", style="normal", variant="normal", weight="normal", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/STIXTwoExtra-BoldItalic.ttf", name="STIXTwoExtra-BoldItalic", style="italic", variant="normal", weight="bold", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/STIXTwoExtra-Bold.ttf", name="STIXTwoExtra-Bold", style="normal", variant="normal", weight="bold", stretch="normal", size="scalable")) = 11.05
findroot: score(GoTfry(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/STIXTwoExtra-BoldItalic.ttf", name="STIXTwoExtra-BoldItalic", style="italic", variant="normal", weight="bold", stretch="normal", size="scalable")) = 10.495
findroot: score(GoTfry(fname="/usr/share/fonts/liberation/LiberationMono-Bold.ttf", name="LiberationMono", style="normal", variant="normal", weight="bold", stretch="normal", size="scalable")) = 10.335
findroot: score(GoTfry(fname="/usr/share/fonts/liberation/LiberationMono-BoldItalic.ttf", name="LiberationMono", style="italic", variant="normal", weight="bold", stretch="normal", size="scalable")) = 11.335
findroot: score(GoTfry(fname="/usr/share/fonts/liberation/LiberationMono-BoldItalic.ttf", name="LiberationMono", style="normal", variant="normal", weight="bold", stretch="normal", size="scalable")) = 10.335
findroot: Matching same score(fname="/home/naeyoon/.local/lib/python3.9/site-packages/matplotlib/mpl-data/fonts/ttf/DeJaVuSans.ttf" with score of 0.950000.

```